



CQF Exam Three

Machine Learning

Bach Van Hoang Bao

June 2023 cohort

FitchLearning

November 9, 2023

1 What are voting classifiers in ensemble learning?

Voting classifiers are based on the idea of aggregating the predictions of multiple classifiers to make a final decision [1]. There are two main types of voting classifiers:

1. Majority Class Labels (Majority/Hard Voting): In majority voting, the predicted class label for a particular sample is the class label that represents the majority (mode) of the class labels predicted by each individual classifier.

E.g., if the prediction for a given sample is

- classifier 1 \rightarrow class 1
- classifier 2 \rightarrow class 1
- classifier 3 \rightarrow class 2

The VotingClassifier (with voting = 'hard') would classify the sample as "class 1" based on the majority class label.

In the cases of a tie, the VotingClassifier will select the class based on the ascending sort order. E.g., in the following scenario

- classifier 1 \rightarrow class 2
- classifier 2 \rightarrow class 1

The class label 1 will be assigned to the sample.

2. Soft Voting Classifier: In contrast to majority voting (hard voting), soft voting returns the class label as argmax of the sum of predicted probabilities.

Specific weights can be assigned to each classifier via the weights parameter. When weights are provided, the predicted class probabilities for each classifier are collected, multiplied by the classifier weight, and averaged. The final class label is then derived from the class label with the highest average probability.

To illustrate this with a simple example, let's assume we have 3 classifiers and a 3-class classification problems where we assign equal weights to all classifiers: $w_1 = 1, w_2 = 1, w_3 = 1$.

The weighted average probabilities for a sample would then be calculated as follows:

Classifier	Class 1	Class 2	Class 3
Classifier 1	$w_1 \times 0.2$	$w_1 \times 0.5$	$w_1 \times 0.3$
Classifier 2	$w_2 \times 0.6$	$w_2 \times 0.3$	$w_2 \times 0.1$
Classifier 3	$w_3 \times 0.3$	$w_3 \times 0.4$	$w_3 \times 0.3$
Weighted Avg	0.37	0.4	0.23

Here, the predicted class label is 2, since it has the highest average probability.

2 Explain the role of the regularization parameter C in a Support Vector Machine (SVM) model. How does varying C affect the model's bias and variance trade-off?

Consider the mathematical equation for the soft margin of non linearly separable data [2] \mathbf{x}_n, y_n :

$$\begin{aligned} \min_{\mathbf{w}, b, \xi} \quad & \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{n=1}^N \xi_n \\ \text{subject to} \quad & y_n(\langle \mathbf{w}, \mathbf{x}_n \rangle + b) \geq 1 - \xi_n \\ & \xi_n \geq 0 \end{aligned}$$

Where:

- \mathbf{w} is the normal vector of the hyper plane
- $\|\mathbf{w}\|^2$ is the regularizer
- \mathbf{x}_n is the feature vector n^{th}
- y_n is the label n^{th}
- ξ_n is the slack term that measures the distance of a positive example x_+ to the positive margin hyperplane ($\langle \mathbf{w}, \mathbf{x} \rangle + b = 1$) when x_+ is on the wrong side.

The parameter $C > 0$ trades off the size of the margin and the total amount of slack that we have. A large value of C implies low regularization, as we give the slack variables larger weight, hence giving more priority to examples that do not lie on the correct side of the margin.

Let take a look at the impact of C in the model using the breast cancer data from Scikit-Learn library.

```
[1]: import pandas as pd
import numpy as np
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.datasets import load_breast_cancer
from sklearn.svm import SVC
import warnings
warnings.filterwarnings("ignore")

cancer = load_breast_cancer()

# The data set is presented in a dictionary form:
print(cancer.keys())
```

```
dict_keys(['data', 'target', 'frame', 'target_names', 'DESCR', 'feature_names',
'filename', 'data_module'])
```

Now we will extract all features into the new data frame and our target features into separate data frames.

```
[2]: df_feat = pd.DataFrame(cancer['data'], columns = cancer['feature_names'])
      # cancer column is our target
      df_target = pd.DataFrame(cancer['target'], columns = ['Cancer'])
      print("Feature Variables: ")
      print(df_feat.info())
```

Feature Variables:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 569 entries, 0 to 568

Data columns (total 30 columns):

#	Column	Non-Null Count	Dtype
0	mean radius	569 non-null	float64
1	mean texture	569 non-null	float64
2	mean perimeter	569 non-null	float64
3	mean area	569 non-null	float64
4	mean smoothness	569 non-null	float64
5	mean compactness	569 non-null	float64
6	mean concavity	569 non-null	float64
7	mean concave points	569 non-null	float64
8	mean symmetry	569 non-null	float64
9	mean fractal dimension	569 non-null	float64
10	radius error	569 non-null	float64
11	texture error	569 non-null	float64
12	perimeter error	569 non-null	float64
13	area error	569 non-null	float64
14	smoothness error	569 non-null	float64
15	compactness error	569 non-null	float64
16	concavity error	569 non-null	float64
17	concave points error	569 non-null	float64
18	symmetry error	569 non-null	float64
19	fractal dimension error	569 non-null	float64
20	worst radius	569 non-null	float64
21	worst texture	569 non-null	float64
22	worst perimeter	569 non-null	float64
23	worst area	569 non-null	float64
24	worst smoothness	569 non-null	float64
25	worst compactness	569 non-null	float64
26	worst concavity	569 non-null	float64
27	worst concave points	569 non-null	float64
28	worst symmetry	569 non-null	float64
29	worst fractal dimension	569 non-null	float64

dtypes: float64(30)

memory usage: 133.5 KB

None

We will split the training data and the test data using 70:30 ratio

```
[3]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(df_feat, np.
    ↪ ravel(df_target),
                                     test_size = 0.30, random_state = 101)
```

And fit the data to our SVC model. After that we can see the hyperparameters of our model using `get_param()` method.

```
[4]: # train the model on train set
model = SVC()
model.fit(X_train, y_train)
```

```
[5]: model.get_params()
```

```
[5]: {'C': 1.0,
      'break_ties': False,
      'cache_size': 200,
      'class_weight': None,
      'coef0': 0.0,
      'decision_function_shape': 'ovr',
      'degree': 3,
      'gamma': 'scale',
      'kernel': 'rbf',
      'max_iter': -1,
      'probability': False,
      'random_state': None,
      'shrinking': True,
      'tol': 0.001,
      'verbose': False}
```

The default value of C is 1. Now we can observe the model performance using the confusion matrix.

```
[6]: # print prediction results
predictions = model.predict(X_test)
print(classification_report(y_test, predictions))
```

	precision	recall	f1-score	support
0	0.95	0.85	0.90	66
1	0.91	0.97	0.94	105
accuracy			0.92	171
macro avg	0.93	0.91	0.92	171
weighted avg	0.93	0.92	0.92	171

Let's change the value of $C = 0.001$ and see the result.

```
[7]: # train the model on train set
model = SVC(C=0.001)
model.fit(X_train, y_train)
# print prediction results
predictions = model.predict(X_test)
print(classification_report(y_test, predictions))
```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	66
1	0.61	1.00	0.76	105
accuracy			0.61	171
macro avg	0.31	0.50	0.38	171
weighted avg	0.38	0.61	0.47	171

2.1 Conclusion

When C is set to a small value (e.g., close to 0), the SVM places a higher emphasis on maximizing the margin and finding the hyperplane that separates the data points with as few errors as possible. In this case, the model is more tolerant of misclassifications (training errors) and is willing to accept a wider margin with a few support vectors. The model's bias is higher, as it tends to underfit the training data by allowing more training errors, but the variance is lower because it maintains a simpler decision boundary.

When C is set to a large value, the SVM imposes a stronger penalty on misclassified points and strives to minimize training errors, even if it means having a narrower margin and more support vectors. The model's bias is lower because it tries to fit the training data as closely as possible, potentially leading to a more complex decision boundary. However, the variance is higher because the model is sensitive to individual data points, which can result in overfitting.

3 Produce a model to predict positive moves (up trend) using machine learning model.

In this section, we will create a Machine Learning (ML) model to predict positive movements using raw financial data from Yahoo Finance. The goal is to emphasize the work flow structure and develop a comprehensive framework from ideation to model evaluation, applying techniques and knowledge from the CQF module 4.

The 7 steps of ML work flow [3]:

Step	Workflow	Remark
1	Ideation	Predict positive moves from the given dataset
2	Data Collection	Download the data from Yahoo Finance and store the data set
3	EDA	Study summary statistics
4	Cleaning Dataset	Trying to resolve the missing data
5	Transformation	Perform feature scaling based on EDA
6	Modelling	Building and training classification model
7	Metrics	Validating the model performance

3.1 STEP 1: Ideation

The objective of the exam is to create a model for predicting upward and downward movements of the underlying asset. This is a classification problem, so we will approach it by assigning a classification label [0] for a downward trend and [1] for an upward trend. We will use the SPY ticker as an example and utilize data from 2008-10-16 to 2023-10-16. The reason for choosing this date range is that it encompasses various financial regimes, including the 2008 financial crash and the 2020 COVID-19 recession.

Given that we are working with financial time series data, the work flow is relatively straightforward. After downloading and storing the data, we will begin by exploring the data to identify any meaningful structures or trends. Next, we'll create the classification labels and address class imbalance. We will apply feature engineering techniques to generate new features from the original data and select the most crucial features for our model. The data will be saved under `../SPY1D.csv`.

Subsequently, we will split the data into training and testing sets, fitting and transforming the training set, and then transforming the test set. This ensures that we avoid any data leakage issues. We will explore the training set to identify trends and significant structures in the data. If necessary, we will scale the data and use the transformed data to train our model.

During the model training process, we will compare the cross-validation accuracy of multiple classification algorithms with default parameters and select the most promising candidate for further tuning. Hyper parameters will be tuned to optimize the selected candidate, leading to the creation of a final model. This final model will be saved as `final_model.joblib` for future use.

In terms of performance measurement, we will employ the confusion matrix and AUC-ROC curve to evaluate our model's performance and document the entire process.

3.2 STEP 2: Data Collection

We will use the `yfinance` package to download daily trading data from Yahoo Finance. The recommended data should span a 5-year period, which is considered sufficient. The downloaded data will be saved in the `.csv` format and can be accessed later using the file name `SPY1D.csv`.

```
[9]: # Download data for TSLA and store as csv file
spy = yf.download("SPY", start = '2008-10-16', end = '2023-10-16' ,
    ↪interval='1D')
spy.to_csv('SPY1D.csv')
```

```
[*****100%*****] 1 of 1 completed
```

```
[10]: spy = pd.read_csv('../module_4/SPY1D.csv')
```

```
[11]: # Verify the downloaded data
spy.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3774 entries, 0 to 3773
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Date        3774 non-null   object
 1   Open        3774 non-null   float64
 2   High        3774 non-null   float64
 3   Low         3774 non-null   float64
 4   Close       3774 non-null   float64
 5   Adj Close   3774 non-null   float64
 6   Volume      3774 non-null   int64
dtypes: float64(5), int64(1), object(1)
memory usage: 206.5+ KB
```

3.3 STEP 3: EDA

Visualize asset path:

```
[12]: import plotly.express as px

fig = px.line(spy, x= 'Date', y='Adj Close', labels = {'Adj Close': 'Close Price_
    ↪(USD)'}, title = 'S&P 500 ETF Trust (SPY) Daily')
fig.show();
```

3.3.1 Calculate returns

We can plot the distribution of returns and the closing price movement to identify any trends or significant information regarding the returns that could be useful.


```
[13]: spy['Returns'] = np.log(spy['Adj Close']).diff()
```



```
[14]: from scipy.stats import norm

# Plot the return histogram
fig = plt.figure(figsize=(15, 7))
ax1 = fig.add_subplot(1, 1, 1)
spy['Returns'].hist(bins=50, ax=ax1)
ax1.set_xlabel('Return')
ax1.set_ylabel('Count')
ax1.set_title('Return distribution')

# Plot the normal distribution
mu = spy['Returns'].mean()
sigma = spy['Returns'].std()
x = np.linspace(mu - 3*sigma, mu + 3*sigma, 100)
plt.plot(x, norm.pdf(x, mu, sigma))
plt.show()
```



The return is definitely not normally distributed. There is a high peak and very fat tails.

3.3.2 Feature Specify

Using the feature list table from the exam, we will generate features based on the historical data we have acquired. Additionally, I've included 10 lagged prices in the feature list, operating on the assumption that historical data may possess predictive capabilities.

```
[15]: # Create features (predictors) list
features_list = []
# Intraday price range
spy['OC'] = spy['Open'] - spy['Close']
```

```

spy['HL'] = spy['High'] - spy['Low']
# Sign of return or momentum
spy['Sign'] = np.sign(spy>Returns)

# Append feature list
features_list.append('OC')
features_list.append('HL')
features_list.append('Sign')

# Pass Returns, Volatility
for r in range(10, 65, 5):
    spy['Ret_'+str(r)] = spy>Returns.rolling(r).sum()
    spy['Std_'+str(r)] = spy>Returns.rolling(r).std()
    features_list.append('Ret_'+str(r))
    features_list.append('Std_'+str(r))

# SMA and EMA
for a in range(20, 200, 10):
    spy['SMA_'+str(r)] = spy['Adj Close'].rolling(r).mean()
    spy['EMA_'+str(a)] = spy['Adj Close'].ewm(span = a).mean()
    features_list.append('SMA_'+str(r))
    features_list.append('EMA_'+str(r))

# Lag price
for lag in range(1, 10):
    spy['lag_' + str(lag)] = spy['Adj Close'].shift(lag)

# Drop NaN values
spy.dropna(inplace=True)

```

3.3.3 Define target

We define the target variable to be whether the ‘SPY’ price will close up or down on the next trading day. If tomorrow’s closing price is greater than today’s closing price by at least 5%, we consider the asset to be “up”; otherwise, it is considered “down.”

We assign a value of 1 to denote an “up” move and 0 to represent a “down” move for the target variable. This target variable can be described as follows:

$$y_t = \begin{cases} 1, & \text{If } p_t < 0.995 \times p_{t+1} \\ 0, & \text{Otherwise} \end{cases}$$

```

[16]: # Define Target
spy['Target'] = np.where(spy['Adj Close'].shift(-1) > 0.995 * spy['Adj_
    ↳Close'],1,0)
# Check output
spy.head(10)

```

I am going to split the data into the `train_set` and `test_set` and perform exploratory data analysis (EDA) and data cleaning exclusively on the `train_set` to prevent any potential data leakage from the EDA process.

```
[17]: # Copy the original data
data = spy.copy().set_index('Date')

[18]: # Specify the features matrix `X`
X = data.drop(['Open', 'Close', 'High', 'Low', 'Adj Close', 'Returns', 'Volume',
→ 'Target'],axis=1)
X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 3714 entries, 2009-01-13 to 2023-10-13
Data columns (total 53 columns):
#   Column      Non-Null Count  Dtype
---  -
0   OC           3714 non-null   float64
1   HL           3714 non-null   float64
2   Sign         3714 non-null   float64
3   Ret_10       3714 non-null   float64
4   Std_10       3714 non-null   float64
5   Ret_15       3714 non-null   float64
6   Std_15       3714 non-null   float64
7   Ret_20       3714 non-null   float64
8   Std_20       3714 non-null   float64
9   Ret_25       3714 non-null   float64
10  Std_25       3714 non-null   float64
11  Ret_30       3714 non-null   float64
12  Std_30       3714 non-null   float64
13  Ret_35       3714 non-null   float64
14  Std_35       3714 non-null   float64
15  Ret_40       3714 non-null   float64
16  Std_40       3714 non-null   float64
17  Ret_45       3714 non-null   float64
18  Std_45       3714 non-null   float64
19  Ret_50       3714 non-null   float64
20  Std_50       3714 non-null   float64
21  Ret_55       3714 non-null   float64
22  Std_55       3714 non-null   float64
23  Ret_60       3714 non-null   float64
24  Std_60       3714 non-null   float64
25  SMA_60       3714 non-null   float64
26  EMA_20       3714 non-null   float64
27  EMA_30       3714 non-null   float64
28  EMA_40       3714 non-null   float64
29  EMA_50       3714 non-null   float64
30  EMA_60       3714 non-null   float64
```

```

31 EMA_70    3714 non-null    float64
32 EMA_80    3714 non-null    float64
33 EMA_90    3714 non-null    float64
34 EMA_100   3714 non-null    float64
35 EMA_110   3714 non-null    float64
36 EMA_120   3714 non-null    float64
37 EMA_130   3714 non-null    float64
38 EMA_140   3714 non-null    float64
39 EMA_150   3714 non-null    float64
40 EMA_160   3714 non-null    float64
41 EMA_170   3714 non-null    float64
42 EMA_180   3714 non-null    float64
43 EMA_190   3714 non-null    float64
44 lag_1     3714 non-null    float64
45 lag_2     3714 non-null    float64
46 lag_3     3714 non-null    float64
47 lag_4     3714 non-null    float64
48 lag_5     3714 non-null    float64
49 lag_6     3714 non-null    float64
50 lag_7     3714 non-null    float64
51 lag_8     3714 non-null    float64
52 lag_9     3714 non-null    float64

```

dtypes: float64(53)

memory usage: 1.5+ MB

```

[19]: # Define label or target vector `y`
      y = data['Target']
      y

```

```

[19]: Date
      2009-01-13    0
      2009-01-14    1
      2009-01-15    1
      2009-01-16    0
      2009-01-20    1
      ..
      2023-10-09    1
      2023-10-10    1
      2023-10-11    0
      2023-10-12    1
      2023-10-13    0
      Name: Target, Length: 3714, dtype: int64

```

```

[20]: # Splitting the datasets into training and testing data.
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪shuffle=False)
      # Output the train and test data size

```

```
print(f"Train and Test Size {len(X_train)}, {len(X_test)}")
```

Train and Test Size 2971, 743

3.3.4 Imbalance class

Since this is a classification problem, it's important to check for any imbalances in our labels.

```
[21]: # class frequency
      c = y_train.value_counts()
      c
```

```
[21]: 1    2361
      0     610
      Name: Target, dtype: int64
```

The label is imbalanced. We will create a weight function and subsequently use it to address our problem when building a model.

```
[22]: # class weight function
      def cwts(label):
          c0, c1 = np.bincount(label)
          w0=(1/c0)*(len(label))/2
          w1=(1/c1)*(len(label))/2
          return {0: w0, 1: w1}
```

```
[23]: # check class weights
      class_weight = cwts(y_train)
      class_weight
```

```
[23]: {0: 2.435245901639344, 1: 0.6291825497670478}
```


3.3.5 Multi collinearity features

Collinear features can adversely affect our model's performance. We will create a function to help us identify and drop these features, and then apply it to our test dataset. Let's also visualize our correlation matrix using the `sns.heatmap()` method.

```
[24]: plt.figure(figsize=(25, 22))
```

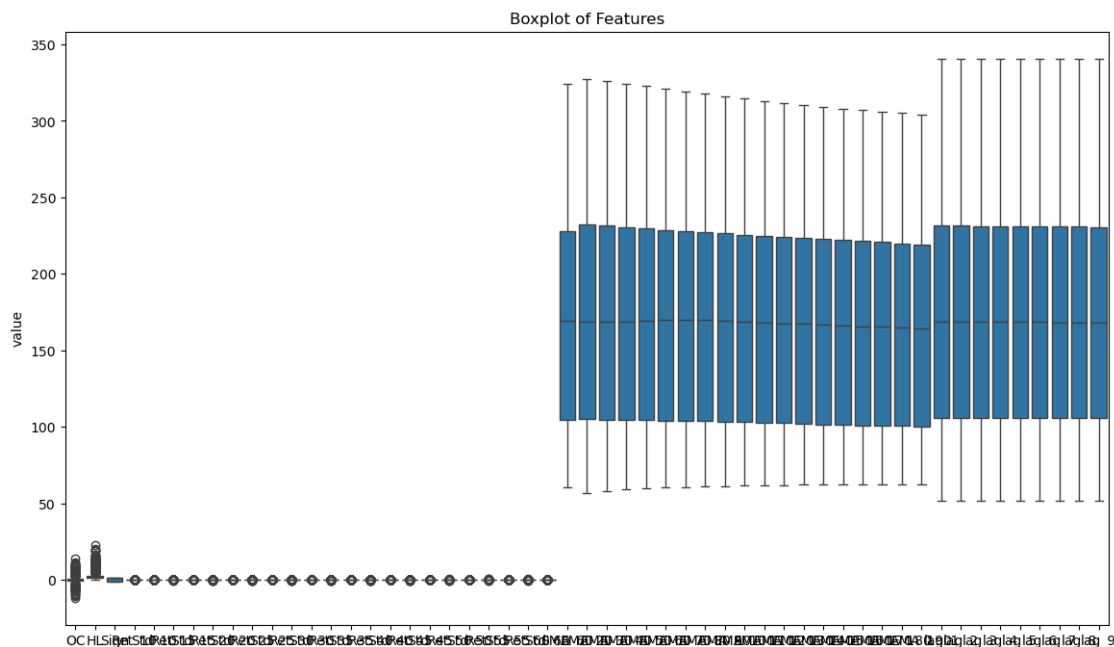
```
# Identify features that are highly correlated
sns.heatmap(X_train.corr().>0.9,
            annot=True,
            annot_kws={"size": 8},
            fmt=".2f",
            linewidth=.5,
            cmap="coolwarm",
            cbar=True); #cmap="crest", viridis, magma
```

```
plt.title('Features Set Correlations');
```



Feature scaling is also a crucial factor in our model's accuracy. We need to scale the data before inputting it into our learning algorithm. We can easily identify features that require scaling by using the `sns.boxplot()` method.

```
[25]: # study the distribution
fig, ax = plt.subplots(figsize=(14,8))
sns.boxplot(x='variable', y='value', data=pd.melt(X_train))
plt.xlabel(' ')
plt.title('Boxplot of Features');
```



Alternatively, we can identify features that require scaling by using the `pd.describe()` method.

```
[26]: X_train.describe()
```

```
[26]:
```

	OC	HL	Sign	Ret_10	Std_10	\
count	2971.000000	2971.000000	2971.000000	2971.000000	2971.000000	
mean	-0.032289	2.180993	0.119488	0.005261	0.009328	
std	1.653484	1.984382	0.991306	0.033085	0.006957	
min	-11.680008	0.299995	-1.000000	-0.265117	0.001264	
25%	-0.750000	1.129997	-1.000000	-0.007946	0.005041	
50%	-0.110001	1.619995	1.000000	0.008294	0.007418	
75%	0.580017	2.490005	1.000000	0.022695	0.011363	
max	13.729996	22.960007	1.000000	0.195407	0.071223	

	Ret_15	Std_15	Ret_20	Std_20	Ret_25	...	\
count	2971.000000	2971.000000	2971.000000	2971.000000	2971.000000	...	

mean	0.007996	0.009473	0.010643	0.009577	0.013273	...
std	0.040366	0.006703	0.045877	0.006543	0.050274	...
min	-0.320822	0.001494	-0.370872	0.002007	-0.409051	...
25%	-0.007578	0.005288	-0.007721	0.005570	-0.006560	...
50%	0.013048	0.007591	0.016520	0.007800	0.019930	...
75%	0.028728	0.011410	0.034884	0.011496	0.040088	...
max	0.241287	0.065627	0.212051	0.059167	0.248100	...

	EMA_190	lag_1	lag_2	lag_3	lag_4	\
count	2971.000000	2971.000000	2971.000000	2971.000000	2971.000000	
mean	163.806877	171.258269	171.175743	171.089990	171.003766	
std	68.167202	71.091955	71.070362	71.040253	71.009352	
min	62.545531	51.386787	51.386787	51.386787	51.386787	
25%	100.199986	105.710148	105.662731	105.651165	105.642242	
50%	164.247036	168.587280	168.572571	168.515594	168.472672	
75%	219.259284	231.435677	231.315384	231.196716	231.115906	
max	304.055049	340.724152	340.724152	340.724152	340.724152	

	lag_5	lag_6	lag_7	lag_8	lag_9
count	2971.000000	2971.000000	2971.000000	2971.000000	2971.000000
mean	170.916197	170.828847	170.742130	170.654513	170.567015
std	70.972773	70.937136	70.902683	70.868631	70.835903
min	51.386787	51.386787	51.386787	51.386787	51.386787
25%	105.641220	105.630589	105.619759	105.619560	105.594227
50%	168.462601	168.411850	168.332382	168.284988	168.247406
75%	231.058403	230.981613	230.868690	230.719635	230.638336
max	340.724152	340.724152	340.724152	340.724152	340.724152

[8 rows x 53 columns]

Some features exhibit significantly higher absolute values compared to the others. For these features, we will use the `MinMaxScaler()` method to scale them appropriately.

3.4 STEP 4: Cleaning Data

From our exploratory data analysis (EDA) process, we have identified multicollinear features. We will develop a function to eliminate these features and then implement it on our training data. Subsequently, we will apply the same function to our test data.

```
[27]: # remove the first feature that is correlated with any other feature
def correlated_features(data, threshold=0.9):
    col_corr = set()
    corr_matrix = X_train.corr()
    for i in range(len(corr_matrix.columns)):
        for j in range(i):
            if abs(corr_matrix.iloc[i, j]) > threshold:
                colname = corr_matrix.columns[i]
                col_corr.add(colname)
    return col_corr
```

```
[28]: # Get the list of remaining features
drop_correlated_features = correlated_features(X_train, threshold=0.9)
```

```
[29]: # drop the highly correlated features
X_train_drop = X_train.drop(drop_correlated_features, axis=1)
X_train_drop.describe()
```

```
[29]:
```

	OC	HL	Sign	Ret_10	Std_10 \
count	2971.000000	2971.000000	2971.000000	2971.000000	2971.000000
mean	-0.032289	2.180993	0.119488	0.005261	0.009328
std	1.653484	1.984382	0.991306	0.033085	0.006957
min	-11.680008	0.299995	-1.000000	-0.265117	0.001264
25%	-0.750000	1.129997	-1.000000	-0.007946	0.005041
50%	-0.110001	1.619995	1.000000	0.008294	0.007418
75%	0.580017	2.490005	1.000000	0.022695	0.011363
max	13.729996	22.960007	1.000000	0.195407	0.071223

	Ret_15	Ret_20	Ret_25	Ret_30	SMA_60
count	2971.000000	2971.000000	2971.000000	2971.000000	2971.000000
mean	0.007996	0.010643	0.013273	0.015933	168.784765
std	0.040366	0.045877	0.050274	0.053629	69.918583
min	-0.320822	-0.370872	-0.409051	-0.392927	60.405024
25%	-0.007578	-0.007721	-0.006560	-0.005451	104.476129
50%	0.013048	0.016520	0.019930	0.021935	169.387304
75%	0.028728	0.034884	0.040088	0.044824	227.957799
max	0.241287	0.212051	0.248100	0.249708	323.832030

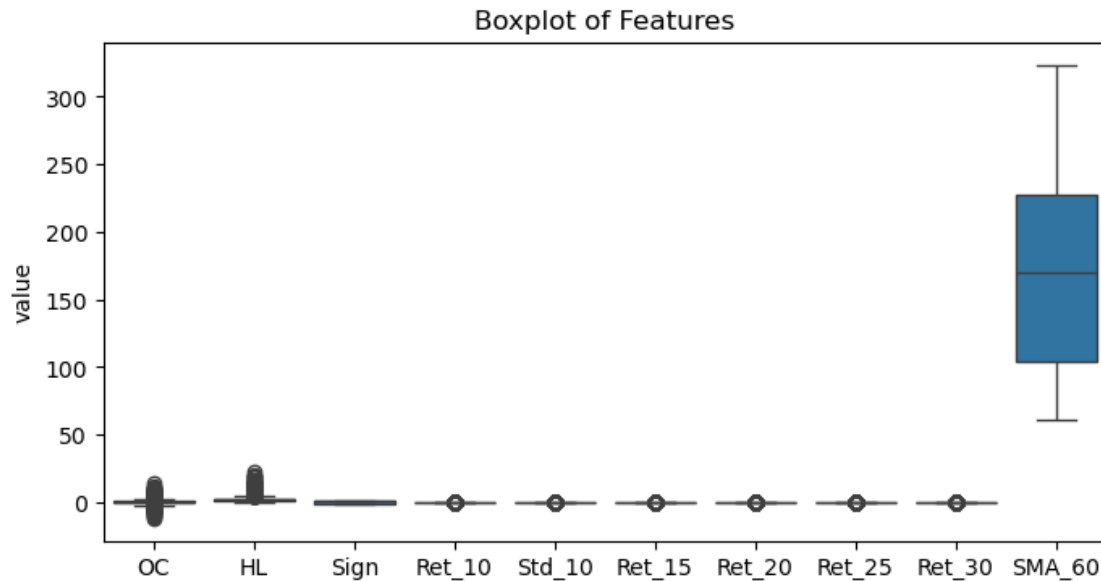
After removing most of the highly correlated features, it appears that past returns, past volatility, SMA (Simple Moving Average), OC (Open-Close), HL (High-Low), and Sign have significant predictive power.

```
[30]: X_test_drop = X_test.drop(drop_correlated_features, axis=1)
```

3.5 STEP 5: Transformation

We will visualize the scale of our data once more before proceeding with feature transformation.

```
[31]: # study the distribution
plt.figure(figsize=(8, 4))
sns.boxplot(x='variable', y='value', data=pd.melt(X_train_drop))
plt.xlabel(' ')
plt.title('Boxplot of Features');
```



The only remaining feature with a high value is SMA_60. We will scale this feature using `MinMaxScaler()`.

```
[32]: minmax = ColumnTransformer([
    ('scaled', MinMaxScaler(), ['SMA_60'])
], remainder = 'passthrough')
```

```
[33]: # Fit and transform the data
sma_60 = minmax.fit_transform(X_train_drop)
```

```
[34]: X_train_dropped_scaled = pd.DataFrame(
    sma_60, columns=minmax.get_feature_names_out(),
    index=X_train_drop.index)
```

```
[35]: X_train_dropped_scaled.describe()
```

```
[35]:
```


	scaled__SMA_60	remainder__OC	remainder__HL	remainder__Sign	\
count	2971.000000	2971.000000	2971.000000	2971.000000	
mean	0.411422	-0.032289	2.180993	0.119488	
std	0.265419	1.653484	1.984382	0.991306	
min	0.000000	-11.680008	0.299995	-1.000000	
25%	0.167299	-0.750000	1.129997	-1.000000	
50%	0.413710	-0.110001	1.619995	1.000000	
75%	0.636050	0.580017	2.490005	1.000000	
max	1.000000	13.729996	22.960007	1.000000	

	remainder__Ret_10	remainder__Std_10	remainder__Ret_15	\
count	2971.000000	2971.000000	2971.000000	
mean	0.005261	0.009328	0.007996	
std	0.033085	0.006957	0.040366	
min	-0.265117	0.001264	-0.320822	
25%	-0.007946	0.005041	-0.007578	
50%	0.008294	0.007418	0.013048	
75%	0.022695	0.011363	0.028728	
max	0.195407	0.071223	0.241287	

	remainder__Ret_20	remainder__Ret_25	remainder__Ret_30
count	2971.000000	2971.000000	2971.000000
mean	0.010643	0.013273	0.015933
std	0.045877	0.050274	0.053629
min	-0.370872	-0.409051	-0.392927
25%	-0.007721	-0.006560	-0.005451
50%	0.016520	0.019930	0.021935
75%	0.034884	0.040088	0.044824
max	0.212051	0.248100	0.249708

Let's visualize our scaled data once more.

```
[36]: fig, ax = plt.subplots(figsize=(18,5))
sns.boxplot(x='variable', y='value', data=pd.melt(X_train_dropped_scaled))
plt.xlabel('After scaled')
plt.title('Boxplot of Features');
```




It appears that the OC and HL columns contain a substantial number of outliers. We will employ the `RobustScaler` to transform these features.

```
[37]: robust = ColumnTransformer([
    ('cols', RobustScaler(), ['OC', 'HL'])
], remainder = 'passthrough')

oc_hl = robust.fit_transform(X_train_drop)

X_train_dropped_scaled = pd.DataFrame(
    oc_hl, columns=robust.get_feature_names_out(),
    index=X_train_drop.index)
```

```
[38]: fig, ax = plt.subplots(figsize=(18,5))
sns.boxplot(x='variable', y='value', data=pd.melt(X_train_dropped_scaled.
↳drop(['remainder__SMA_60'],axis=1)))
plt.xlabel('After scaled')
plt.title('Boxplot of Features');
```



Now we can construct a preprocessing transformer that applies the specified transformations to particular columns. We will fit and transform the training data and subsequently transform the test data.

```
[39]: # Instantiate transformer
preprocessing = ColumnTransformer([
    ('MinMax', MinMaxScaler(), ['SMA_60']),
    ('Robust', RobustScaler(), ['OC', 'HL'])
], remainder = 'passthrough')

[40]: # Fit and transform train set
train_transformed = preprocessing.fit_transform(X_train_drop)
X_train_transformed = pd.DataFrame(
    train_transformed, columns=preprocessing.get_feature_names_out(),
    index=X_train_drop.index)

# Transform test set
test_transformed = preprocessing.transform(X_test_drop)
X_test_transformed = pd.DataFrame(
    test_transformed, columns=preprocessing.get_feature_names_out(),
    index=X_test_drop.index)
```

3.6 STEP 6: Modelling

We will compare the default settings of some classifiers using the cross-validation technique to identify potential candidates for our final model. Additionally, we will use the `class_weight` parameter to address the previously identified class imbalance problem.

```
[41]: # cross-validation
tscv = TimeSeriesSplit(n_splits=5)

[42]: # specify estimators
random_state = 42
dtc = DecisionTreeClassifier(class_weight=class_weight)
rfc = RandomForestClassifier(max_depth = 5 ,class_weight=class_weight,
    ↪random_state=random_state)
knn = KNeighborsClassifier()
gbc = GradientBoostingClassifier(random_state=random_state)
svc = SVC(class_weight=class_weight, random_state=random_state)

[43]: # get cv scores
clf = [dtc, rfc, knn, gbc, svc]
for estimator in clf:
    score = cross_val_score(estimator, X_train_transformed, y_train, scoring =
    ↪'accuracy', cv=tscv, n_jobs=-1)
    print(f"The accuracy score of {estimator} is: {score.mean():0.4}")
```

The accuracy score of DecisionTreeClassifier(class_weight={0: 2.435245901639344, 1: 0.6291825497670478}) is: 0.6505

The accuracy score of RandomForestClassifier(class_weight={0: 2.435245901639344, 1: 0.6291825497670478}, max_depth=5, random_state=42) is: 0.6949

The accuracy score of KNeighborsClassifier() is: 0.7483

The accuracy score of GradientBoostingClassifier(random_state=42) is: 0.5192

The accuracy score of SVC(class_weight={0: 2.435245901639344, 1: 0.6291825497670478}, random_state=42) is: 0.5107

It appears that the RandomForestClassifier() and KNeighborsClassifier() have the highest scores. Given that the KNeighborsClassifier() may not perform well with imbalanced classes, we will concentrate on building the model using the RandomForestClassifier().

3.6.1 Base Model

The default values for the parameters that determine the size of the trees (e.g., `max_depth`, `min_samples_leaf`, etc.) result in fully grown and unpruned trees, which have the potential to overfit our model. To address this, I will set `max_depth` to 5 and then fine-tune this hyperparameter later.

```
[44]: base_model = RandomForestClassifier(max_depth = 5, class_weight=class_weight,
    ↪random_state=random_state)
base_model.fit(X_train_transformed,y_train)
```

```
print (classification_report(y_train[-252:], base_model.  
    ↪predict(X_train_transformed[-252:])))
```

	precision	recall	f1-score	support
0	0.48	0.88	0.62	67
1	0.94	0.66	0.77	185
accuracy			0.72	252
macro avg	0.71	0.77	0.70	252
weighted avg	0.82	0.72	0.73	252

3.6.2 Tuning Hyper-params

We will obtain all the parameters and define our hyperparameter grid.

```
[45]: model = RandomForestClassifier(class_weight=class_weight,
    ↪random_state=random_state, n_jobs=-1)
```

```
[46]: model.get_params()
```

```
[46]: {'bootstrap': True,
      'ccp_alpha': 0.0,
      'class_weight': {0: 2.435245901639344, 1: 0.6291825497670478},
      'criterion': 'gini',
      'max_depth': None,
      'max_features': 'sqrt',
      'max_leaf_nodes': None,
      'max_samples': None,
      'min_impurity_decrease': 0.0,
      'min_samples_leaf': 1,
      'min_samples_split': 2,
      'min_weight_fraction_leaf': 0.0,
      'n_estimators': 100,
      'n_jobs': -1,
      'oob_score': False,
      'random_state': 42,
      'verbose': 0,
      'warm_start': False}
```

As mentioned earlier, we will include `max_depth`, `max_leaf_nodes`, and `n_estimators` in our hyperparameter grid for tuning to prevent overfitting. Additionally, since we are dealing with an imbalanced classification problem, we will experiment with different loss functions to determine their impact on model performance during the hyperparameter search.

```
[47]: # Hyper parameter optimization
param_grid = { 'criterion': ['gini', 'entropy', 'log_loss'],
               'max_depth': [80, 90, 100, 110],
               'max_features': [2, 3],
               'min_samples_leaf': [3, 4, 5],
               'min_samples_split': [8, 10, 12],
               'n_estimators': [100, 200, 300, 1000]
             }

[48]: # perform random search
gs = GridSearchCV(model, param_grid, scoring='f1', cv=tscv, verbose=0, n_jobs=-1)
gs.fit(X_train_transformed, y_train)

[48]: GridSearchCV(cv=TimeSeriesSplit(gap=0, max_train_size=None, n_splits=5,
test_size=None),
                  estimator=RandomForestClassifier(class_weight={0:
2.435245901639344,
                                                    1:
0.6291825497670478},
                                                    n_jobs=-1, random_state=42),
                  n_jobs=-1,
                  param_grid={'criterion': ['gini', 'entropy', 'log_loss'],
                              'max_depth': [80, 90, 100, 110],
                              'max_features': [2, 3], 'min_samples_leaf': [3, 4, 5],
                              'min_samples_split': [8, 10, 12],
                              'n_estimators': [100, 200, 300, 1000]},
                  scoring='f1')

[49]: # best parameters
gs.best_params_

[49]: {'criterion': 'entropy',
      'max_depth': 80,
      'max_features': 2,
      'min_samples_leaf': 3,
      'min_samples_split': 8,
      'n_estimators': 1000}

[50]: # best score
gs.best_score_

[50]: 0.8830577734916052
```


3.7 STEP 7: Metrics

After fine-tuning our model and conducting a search for the best hyperparameters, we will evaluate our model's performance and compare it to our base model.

```
[51]: # Refit the XGB Classifier with the best params
final_model = RandomForestClassifier(class_weight=class_weight,
    ↪random_state=random_state, n_jobs=-1, **gs.best_params_)
final_model.fit(X_train_transformed, y_train)

[51]: RandomForestClassifier(class_weight={0: 2.435245901639344,
    1: 0.6291825497670478},
    criterion='entropy', max_depth=80, max_features=2,
    min_samples_leaf=3, min_samples_split=8,
    n_estimators=1000, n_jobs=-1, random_state=42)

[52]: # Predicting the test dataset
y_pred = final_model.predict(X_test_transformed)
# Measure Accuracy
acc_train = accuracy_score(y_train, final_model.predict(X_train_transformed))
acc_test = accuracy_score(y_test, y_pred)
# Print Accuracy
print(f'\n Training Accuracy \t: {acc_train :0.4} \n Test Accuracy \t\t:
    ↪{acc_test :0.4}')
```

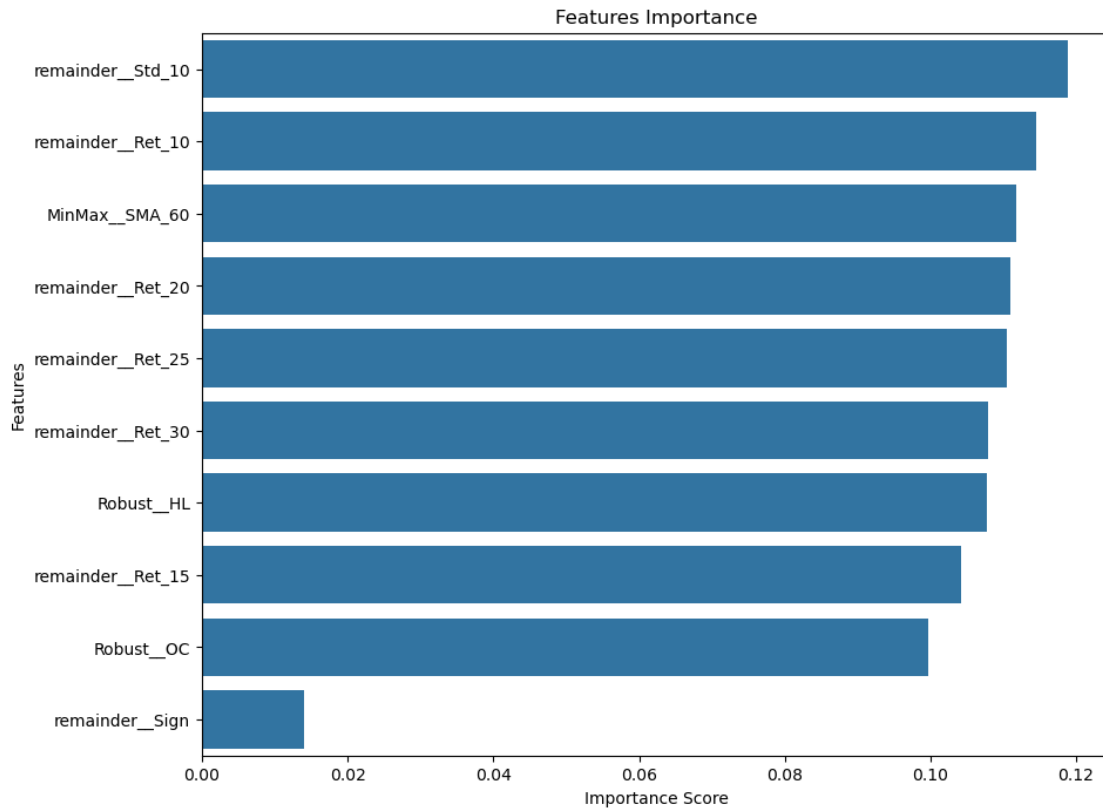
```
Training Accuracy      : 0.9973
Test Accuracy          : 0.6824
```

Our final model outperforms the base model, but it appears to suffer from severe overfitting.

```
[53]: # Cross validation score
score = cross_val_score(final_model, X_train_transformed, y_train, cv=tscv)
print(f'Mean CV Score : {score.mean():0.4}')
```

```
Mean CV Score : 0.7947
```

```
[54]: # Plot feature importance
fig, ax = plt.subplots(figsize=(10,8))
feature_imp = pd.DataFrame({'Importance Score': final_model.
    ↪feature_importances_, 'Features': X_train_transformed.columns}).
    ↪sort_values(by='Importance Score', ascending=False)
sns.barplot(x=feature_imp['Importance Score'], y=feature_imp['Features'])
ax.set_title('Features Importance');
```



It appears that the most important features are the asset returns, volatilities, and H-L (High-Low) values. There is some predictive power in the SMA_60 feature, while the Sign feature contributes almost no predictive power.

```
[55]: # Classification Report
print(classification_report(y_test, y_pred))
```


	precision	recall	f1-score	support
0	0.32	0.15	0.21	200
1	0.74	0.88	0.80	543
accuracy			0.68	743
macro avg	0.53	0.52	0.50	743
weighted avg	0.62	0.68	0.64	743

```
[56]: # Display confusion matrix
disp = ConfusionMatrixDisplay.from_estimator(
    final_model,
    X_test_transformed,
    y_test,
```

```

        display_labels=final_model.classes_,
        cmap=plt.cm.Blues
    )
disp_ax_.set_title('Final Model')
plt.show()

```




Our model is performing well when predicting the majority class but struggles when predicting the minority class.

```

[57]: # Display ROC Curve
disp_roc = RocCurveDisplay.from_estimator(
    final_model,
    X_test_transformed,
    y_test,
    name='Tuned Random Forest')
disp_roc.ax_.set_title('ROC Curve')
plt.plot([0,1], [0,1], linestyle='--')
plt.show()

```



When analyzing the ROC curve, it becomes evident that our model's performance is only marginally better than random chance. This suggests that the model may not effectively discriminate between positive and negative outcomes. To enhance its predictive power and better address the minority class, we may need to further refine our model or consider additional strategies such as resampling techniques or employing different algorithms.

```
[58]: # Saving final model
      from joblib import dump, load
      dump(clf, 'final_model.joblib')
```

```
[58]: ['final_model.joblib']
```

3.8 Conclusion

We have successfully created a machine learning model capable of predicting upward and downward movements of the underlying asset using raw data obtained from Yahoo Finance APIs. Our process involved feature engineering to create a relevant feature set, feature selection to determine the most important features, and data transformation for our chosen set of features. The Random Forest Classifier is used as a candidate, fine-tuned, and optimized as our final model.

However, there are certain limitations to our model. The imbalanced nature of our labels, particularly in the context of financial time series, poses a challenge, which we've partly addressed using the `class_weight` function. The daily price data is limited, providing relatively weak data structure for our algorithm; exploring higher frequency data may lead to a more robust training set. Additionally, our model exhibits overfitting, suggesting the potential need for a more complex model or better feature engineering techniques.

4 References

- [1] [Scikit-learn 1.3.2 - Voting Classifier](#)
- [2] Mathematics for Machine Learning - Marc Peter Deisenroth - p.380
- [3] Introduction to ML using Scikit-learn - Pythonlab 09